

Dessi Georgieva

AI Systems Engineer · Applied AI, Agents and Product Platforms

London, UK · dessi.georgieva8@gmail.com · [LinkedIn](#) · [GitHub](#) · [Application evidence](#)

Profile

I build AI systems that use tools, complete multi-step work, and recover when something fails. I decide what an agent may know and do, how its work is checked, and when a person should review it. My experience spans open source products and enterprise systems. I am applying for OpenAI roles in London across applied AI, Codex agents, ads solutions, privacy, and integrity.

Selected Product and Engineering Work

Founder platform and product architecture

Human Systems Platform · Evidence-backed learning and professional discovery · [First public product](#)

- I am developing Human Systems Platform for people and organisations that need credible evidence of capability in AI-mediated work. It turns selected experience into private learning, owner-approved public evidence, and testable hypotheses about where that capability may create value.
- I designed it as a federated product: Dessi Space preserves private experience, Learning Foundry supports learning, and DG-OS publishes approved profiles. I have built the shared review, approval, protected persistence, learning-evidence, and publication layers. Product integration and commercial validation are in progress.

Agent execution, policy, and recovery

Agentic Commerce Control Plane · Governed optimisation and learning for commerce · [Public repository](#)

- I built Agentic Commerce to help brands test and improve how they describe products for paid placements and organic agent-led discovery, without giving an AI system unchecked control over business actions.
- I designed a Bayesian-style learning loop that updates brand- and product-scoped beliefs as evidence arrives and carries supported patterns into later query and copy generation. It recommends whether to promote, revise, or reject each variant, while the supervised runtime keeps tools, evidence, memory, approvals, and recovery under explicit control.

Identity, delegated authority, and bounded execution

Gateplane Agent Control Plane · Enterprise control for agent work · [Product overview](#)

- I built Gateplane for organisations that need AI agents to work across users, data, and tools without weakening enterprise access controls.
- I tied identity, delegated authority, task-bound approvals, isolated workspaces, tool scope, evidence, and publication into one control plane around model-proposed work.

Codex capabilities, evaluation, and provenance

Learning Foundry · Learning with AI without confusing output with understanding · [Submitted repository state](#)

- I built Learning Foundry for people who want to learn with AI without treating an agent's successful output as proof of their own understanding. It supports explanation, practice, application, evaluation, and revision while separately developing agent capabilities.
- I built the append-only evidence ledger, deterministic projections, learning checks, capability lifecycle, and consent-gated Codex adapter. Within Human Systems Platform, it contributes curated learning evidence while its private ledger remains inside its own boundary.

Context construction, provider reliability, and agent UX

DG-OS Agent Runtime · Owner-approved public profiles and discovery · [First live profile](#) · [Repository](#)

- I built DG-OS as a public profile and discovery product where visitors can inspect a person's systems, evidence, writing, and development or ask questions grounded in reviewed sources. My profile is its first live public instance.
- Within Human Systems Platform, I use DG-OS to publish approved profiles for work, collaboration, research, or investment conversations. The runtime combines deterministic navigation, evidence retrieval, provider isolation, streaming, and source-aware diagnostics.

Professional Experience

Engineer

Performics Innovations Lab • Publicis Media • London • Nov 2023 - Present

- I design and build AI and data systems across backend services, cloud workflows, multi-tenant platform controls, behavioural modelling, and marketing-intelligence interfaces.
- I build the Programmatic plugin and agent harness across skill selection, a typed interface and backend tools, tenant-bound execution policy, human approval, evidence handling, and recovery.
- I delivered backend services and cloud data workflows for an award-recognised ecommerce optimisation platform.
- I describe employer work only at responsibility and outcome level. Private code, clients, logs, and infrastructure remain confidential.

Senior Business Intelligence Analyst

Publicis Media • London • Mar 2023 - Nov 2023

- I bridged analytics and decision systems, turning recurring analysis into reusable services and platformised workflows.

Business Intelligence Manager

Jellyfish • London • Jan 2021 - Mar 2023

- I built analytics applications and data workflows spanning interactive products, cloud ETL, and data-lake patterns.

Data Consultant / SQL Developer / Data Analyst

Selected contracts • London • 2017 - 2020

- I delivered CRM, analytics, segmentation, automation, and experimentation systems across agencies, startups, and media organisations.

AI Engineering Capabilities

- **Harness architecture:** Design the environment around a model, including context, memory, tools, instructions, state, budgets, and interaction with surrounding systems.
- **Agent legibility:** Turn product goals, repository knowledge, interfaces, and feedback into structures that agents can inspect and follow.
- **Deterministic control:** Keep identity, permissions, approvals, tool scope, evidence, and consequential effects outside model authority.
- **Evaluation and improvement:** Test model behaviour, tool use, system interactions, and end results, then convert failures and feedback into reviewable changes.
- **Reliability and recovery:** Build explicit execution states, traces, receipts, retries, replay, fallback, cancellation, and recovery.
- **Human attention:** Decide when an agent should continue alone, ask for clarification, request approval, or surface uncertainty.

Evidence Boundaries

- Public repositories demonstrate agent architecture, evaluation discipline, and product implementation; they do not establish production-scale Codex traffic.
- I do not claim model fine-tuning experience or a public coding-task solve-rate benchmark.
- Token cost, latency, capacity, and employer-system operational measurements are not publicly available.
- Current learning changes evidence, beliefs, memory, policy, and versioned harness configuration. It does not silently retrain or self-modify a foundation model.

Education

MA Applied Human Rights, University of York • 2009 - 2011

BA Philosophy (Philosophy of Science), Sofia University • 2003 - 2007