

Dessi Georgieva

AI Systems Engineer · Applied AI, Agents and Product Platforms

London, UK · dessi.georgieva8@gmail.com · [LinkedIn](#) · [GitHub](#) · [Application evidence](#)

Profile

I build governed AI products that use tools, complete multi-step work, and recover when something fails. I decide what an agent may know and do, how its work is checked, and when a person should review it. My experience spans enterprise systems, open-source products, and Human Systems Platform, my founder product for evidence-led human-AI learning and opportunity research. I am applying for OpenAI roles in London across applied AI, Codex agents, ads solutions, privacy, and integrity.

Selected Product and Engineering Work

Founder platform and product architecture

Human Systems Platform · Evidence-backed learning, opportunity research, and professional discovery · [First public product](#)

- I am developing Human Systems Platform for people and organisations that need a more credible way to develop capability, represent it through evidence, and discover where it may create value. Dessi Space preserves private continuity; Learning Foundry develops human and agent capability; Opportunity Studio researches opportunities and markets; DG-OS publishes approved evidence; Organization Foundry remains planned.
- I have built review, approval, protected persistence, learning-evidence, publication, and Opportunity Studio preview components. The preview validates approved missions and returns cited hypotheses with visible coverage; research remains manual, and it cannot act for the owner.
- Learning Foundry (MIT) and DG-OS (AGPL-3.0-only) are open source and adaptable. Human Systems Platform, Dessi Space, and Opportunity Studio remain private, keeping personal evidence outside public code.

Agent execution, policy, and recovery

Agentic Commerce Control Plane · Governed optimisation and learning for commerce · [Public repository](#)

- I built Agentic Commerce to test and improve product messaging for paid placements and agent-led discovery without giving AI unchecked control over business actions.
- I designed a Bayesian-style loop that updates brand- and product-scoped beliefs, recommends whether to promote, revise, or reject a variant, and keeps tools, evidence, memory, approvals, and recovery under explicit control.

Identity, delegated authority, and bounded execution

Gateplane Agent Control Plane · Enterprise control for agent work · [Product overview](#)

- I built Gateplane for organisations that need AI agents to work across users, data, and tools without weakening enterprise access controls.
- I tied identity, delegated authority, task-bound approvals, isolated workspaces, tool scope, evidence, and publication into one control plane around model-proposed work.

Codex capabilities, evaluation, and provenance

Learning Foundry · Open-source learning with AI without confusing output with understanding · [MIT-licensed repository](#)

- I built Learning Foundry's append-only evidence ledger, deterministic projections, learning checks, capability lifecycle, and consent-gated Codex adapter. It develops human understanding and agent capabilities separately. The MIT-licensed implementation began during OpenAI Build Week and now continues as Human Systems Platform's general learning module.

Context construction, provider reliability, and agent UX

DG-OS Agent Runtime · Open-source public profiles and discovery · [First live profile](#) · [AGPL-3.0-only repository](#)

- I built DG-OS as a reusable public profile and discovery product with inspectable systems, evidence, writing, and source-grounded questions. Within Human Systems Platform, it is the public expression module; my profile is its first live instance.
- The runtime combines deterministic navigation, evidence retrieval, provider isolation, streaming, and source-aware diagnostics.

Professional Experience

Engineer

Performics Innovations Lab · Publicis Media · London · Nov 2023 - Present

- I design and build AI and data systems across backend services, cloud workflows, multi-tenant platform controls, behavioural modelling, and marketing-intelligence interfaces.
- I build the Programmatic plugin and agent harness across skill selection, a typed interface and backend tools, tenant-bound execution policy, human approval, evidence handling, and recovery.
- I delivered backend services and cloud data workflows for an award-recognised ecommerce optimisation platform.
- I describe employer work only at responsibility and outcome level. Private code, clients, logs, and infrastructure remain confidential.

Senior Business Intelligence Analyst

Publicis Media · London · Mar 2023 - Nov 2023

- I bridged analytics and decision systems, turning recurring analysis into reusable services and platformised workflows.

Business Intelligence Manager

Jellyfish · London · Jan 2021 - Mar 2023

- I built analytics applications and data workflows spanning interactive products, cloud ETL, and data-lake patterns.

Data Consultant / SQL Developer / Data Analyst

Selected contracts · London · 2017 - 2020

- I delivered CRM, analytics, segmentation, automation, and experimentation systems across agencies, startups, and media organisations.

AI Engineering Capabilities

- **Harness architecture:** Design the environment around a model, including context, memory, tools, instructions, state, budgets, and interaction with surrounding systems.
- **Agent legibility:** Turn product goals, repository knowledge, interfaces, and feedback into structures that agents can inspect and follow.
- **Deterministic control:** Keep identity, permissions, approvals, tool scope, evidence, and consequential effects outside model authority.
- **Evaluation and improvement:** Test model behaviour, tool use, system interactions, and end results, then convert failures and feedback into reviewable changes.
- **Reliability and recovery:** Build explicit execution states, traces, receipts, retries, replay, fallback, cancellation, and recovery.
- **Human attention:** Decide when an agent should continue alone, ask for clarification, request approval, or surface uncertainty.

Evidence Boundaries

- Public repositories demonstrate agent architecture, evaluation discipline, and product implementation; they do not establish production-scale Codex traffic.
- I do not claim model fine-tuning experience or a public coding-task solve-rate benchmark.
- Token cost, latency, capacity, and employer-system operational measurements are not publicly available.
- Current learning changes evidence, beliefs, memory, policy, and versioned harness configuration. It does not silently retrain or self-modify a foundation model.

Education

MA Applied Human Rights, University of York · 2009 - 2011

BA Philosophy (Philosophy of Science), Sofia University · 2003 - 2007